

The Predictability of Cardiovascular Disease in Patients

Allison Beatty(yxn127),
Joshua Ching (mrf145),
Ruyi Liang (hua075),
Ian Scarff (iie728)

Abstract

There are many risk factors, other than cholesterol and blood pressure, that play an important role in determining whether or not a patient is developing cardiovascular disease. The dataset contains information describing different aspects of patient characteristics, collected at the time of medical examination. The goal of the study is to build a good model for predicting whether a patient has cardiovascular disease, with a high degree of predictive accuracy. The dataset contains 70,000 observations and 12 variables (5 continuous and 6 categorical) describing characteristics for each patient.

Background

Cardiovascular diseases (CVDs) are a group of disorders of the heart and blood vessels such as: heart attack, stroke, heart failure, arrhythmia, and other heart diseases. CVDs are the leading cause of death and a major cause of disability worldwide. “An estimated 17.9 million people died from CVDs in 2016, representing 31% of all global deaths. Of these deaths, 85% are due to heart attack and stroke.”[5] Cardiovascular diseases are also the leading cause of death for men, women, and people of most racial and ethnic groups in the United States. “More than 859,000 Americans die of heart disease, stroke, or other cardiovascular diseases every year—that’s one-third of all US deaths.”[4]

In terms of the national economy, cardiovascular diseases take a financial toll, “costing \$213.8 billion a year to the healthcare system and causing \$137.4 billion in lost productivity from premature death alone.”[4] The U.S. government, primarily the Centers for Disease Control and Prevention (CDC), have responded to such national health concerns by initiating various health studies, branches, and prevention programs in the support of public health efforts that address cardiovascular diseases. For example, “the CDC’s Division for Heart Disease and Stroke Prevention (DHDSP) works with partners across government, public health, health care, and private sectors to improve prevention, detection, and control of heart disease and stroke risk factors, with a focus on high blood pressure and high cholesterol.”[4]

“The aging population, obesity epidemic, underuse of prevention strategies, and suboptimal control of risk factors could exacerbate the future CVD burden.”[2] With such a pressing national health issue that is ongoing today, it is not surprising that business will arise in predicting the significant risk factors associated with those diseases, for the needs in further understanding of the underlying relationship. This would assist physicians in the early detection and management of cardiovascular diseases.

Variable Introduction and Definitions

This dataset contains 70,000 records of patient data and contains 11 predictive features, plus the target variable “Disease,” with no missing values. The dataset can be found at <https://www.kaggle.com/sulianova/cardiovascular-disease-dataset>. Below is a list of the variables as found in the dataset with their descriptions:

Variable Name	Description	Data Type - Unit/Levels
Age	Patient’s age.	Continuous - # of days.
Height	Patient’s height.	Continuous - centimeters.
Weight	Patient’s weight.	Continuous - kilograms.
Gender	Patient’s gender.	2 level Categorical: 1 = woman, 2 = male.
Systolic	Patient’s systolic blood pressure.	Continuous - millimeters of mercury (mmHg).
Diastolic	Patient’s diastolic blood pressure.	Continuous - millimeters of mercury (mmHg).
Cholesterol	Patient’s cholesterol level.	3 level Categorical: 1 = normal, 2 = above normal, 3 = well above normal.
Glucose	Patient’s glucose level.	3 level Categorical: 1 = normal, 2 = above normal, 3 = well above normal.
Smoking	Indicator for whether a patient smokes.	2 level Categorical: 0 = No, 1 = Yes.
Alcohol	Indicator for whether a patient consumes alcohol.	2 level Categorical: 0 = No, 1 = Yes.
Physical	Indicator for whether a patient performs physical activity.	2 level Categorical: 0 = No, 1 = Yes.
Disease	Indicator for whether a patient has a CVD.	2 level Categorical: 0 = No, 1 = Yes.


The following will analyze the relationship between the predictor variables and the response variable, “Disease”. The data will be preprocessed, which includes tasks such as transformation of continuous variables, as well as removal of highly correlated predictors. Both linear models and nonlinear models will be used to see which models perform well with the training data as well as examine the predictability for the testing data.

Data Preprocessing, Exploration, and Transformations

Many statistical models require data to be preprocessed prior to modeling. Preprocessing the data not only helps us appropriately tune the model and improve the predictive performance of the model, but it also decreases the computational time and complexity of the model. The first step we took was to check our data for any missing values. We found that our dataset did not contain any missing values.

We then adjusted the units and levels used in the original data. For “Age,” we converted from the number of days to the number of years (accounting for leap years by dividing by 365.25). We next converted “Weight” and “Height” from metric to imperial. Lastly, we adjusted all categorical levels that do not start at zero to start at zero, so that we can establish a baseline.

The next step we took was to visualize the data to see if there were any other problems. First, we examined the continuous variables (Age, Height, Weight, Systolic, and Diastolic). Below are boxplots of the continuous variables. Looking at these plots, we noticed that there are many illogical/unnatural values.



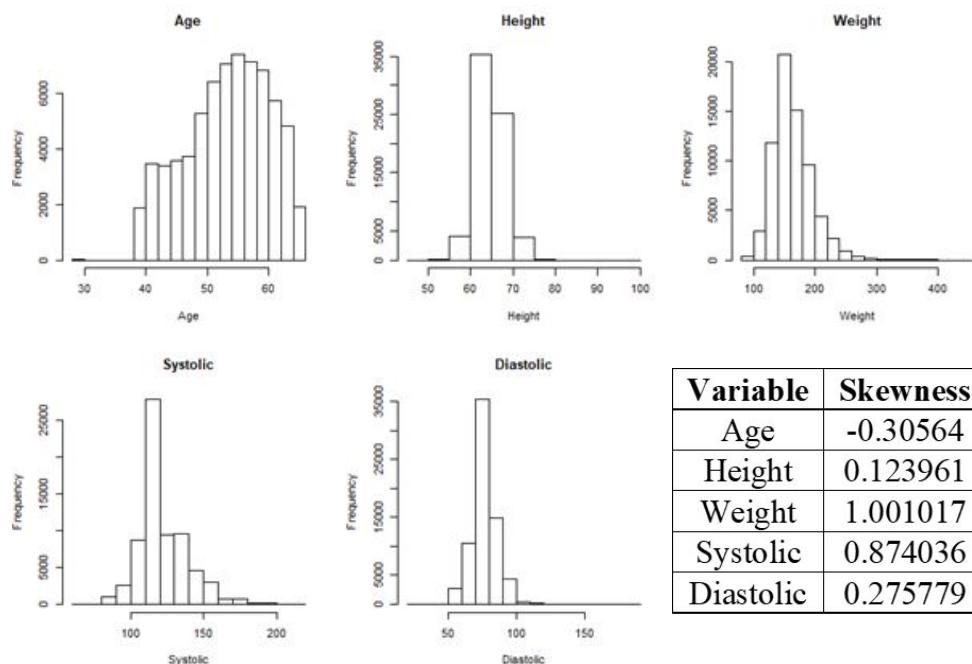
According to the figure displayed above, there are some blood pressures that are higher and lower than what is physically possible or that are not already a medical emergency. In addition, we found that there were some “Diastolic” values that were greater than “Systolic” values for a given observation. This is physically impossible. We can also see that “Weight” has some very low illogical values, given that the lowest value for “Age” is around 29. This same argument can also be applied to very below-average values for “Height.”

There were many assumptions we could have made to try to adjust these illogical values back to what is considered logical, but we decided to focus our study on observations that were already logical. To do this, we filtered the data based on the following restrictions:

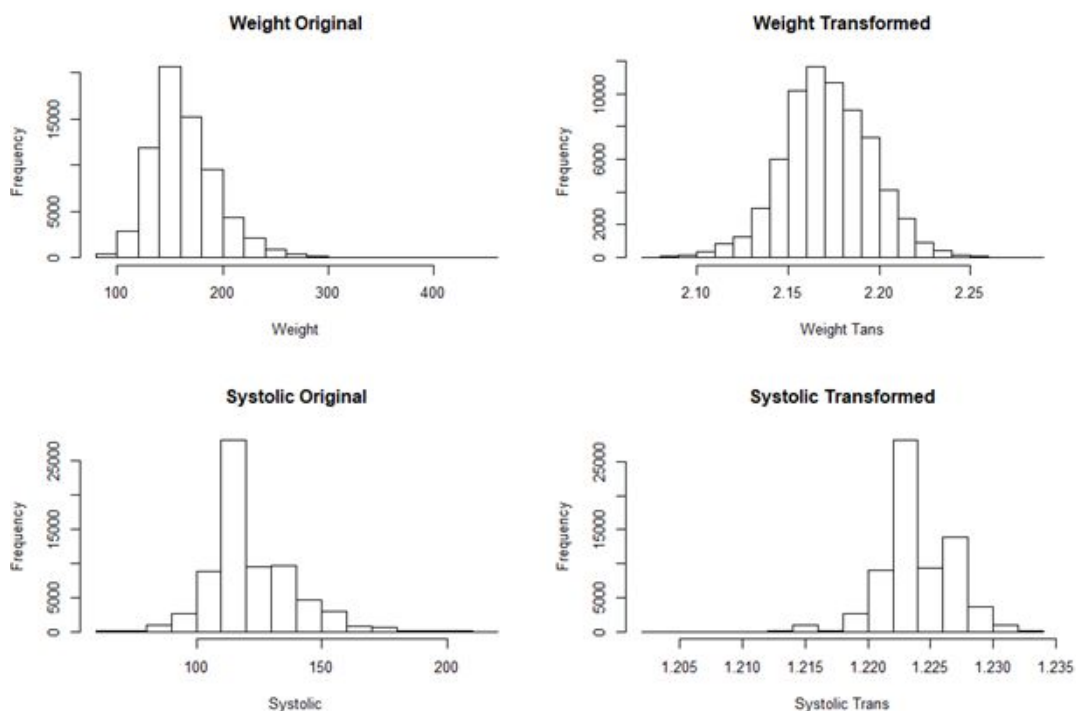
1. Blood pressure can range from hypotension to hypertension. These include “Systolic” values ranging from 50 mmHg to 220 mmHg and “Diastolic” values ranging from 20 mmHg to 190 mmHg.
2. “Diastolic” values must be less than “Systolic” values.
3. The lowest “Weight” value is 80 pounds.
4. The lowest “Height” value is 48 inches.

Even with these restrictions, only approximately 2.05% of the original data was removed and 68,559 observations remain. This restricted data was then used for the remainder of the project.

With this new data, we next examined the distributions of the continuous variables and possible transformations. The figures displayed below are histograms of each of the five continuous variables (Age, Height, Weight, Systolic, and Diastolic) and their measures of skewness.



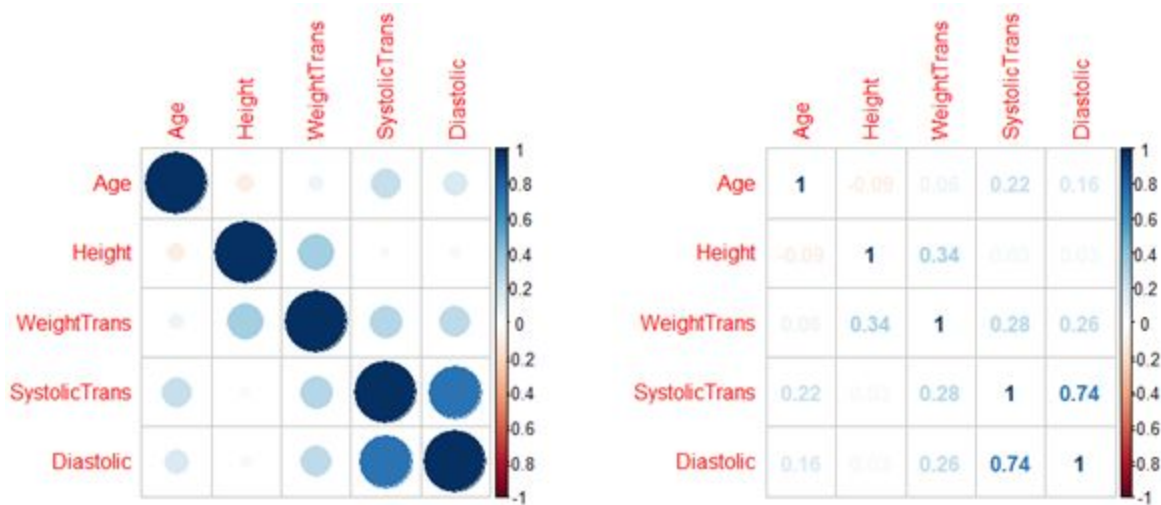
From these plots and tables, we can see that each variable has different levels of skewness. For this project, we considered predictors that have a skewness value between -0.5 and 0.5 to be nearly symmetric, moderately skewed if the absolute value is between 0.5 and 1, and heavily skewed otherwise. Based on this, we examined possible transformations for “Systolic” and “Weight.” Using the Box-Cox transformation method, optimal lambda values for “Systolic” and “Weight” were approximately -0.5 and -1, respectively. The plots displayed below compare the histograms between the original variables and the transformed variables. In addition, the table below compares the original skewness value to the new skewness value for each variable.



Variable	Weight	WeightTrans	Systolic	SystolicTans
Skewness	1.001017	-0.00207732	0.874036	-0.0429108

According to the histograms and table displayed above, we can see that these transformations significantly reduced the skewness in each of these variables. In addition to Box-Cox, we examined other data transformation methods, such as scaling, centering, and combinations of all three methods, but found that using Box-Cox alone provided the best skewness values.

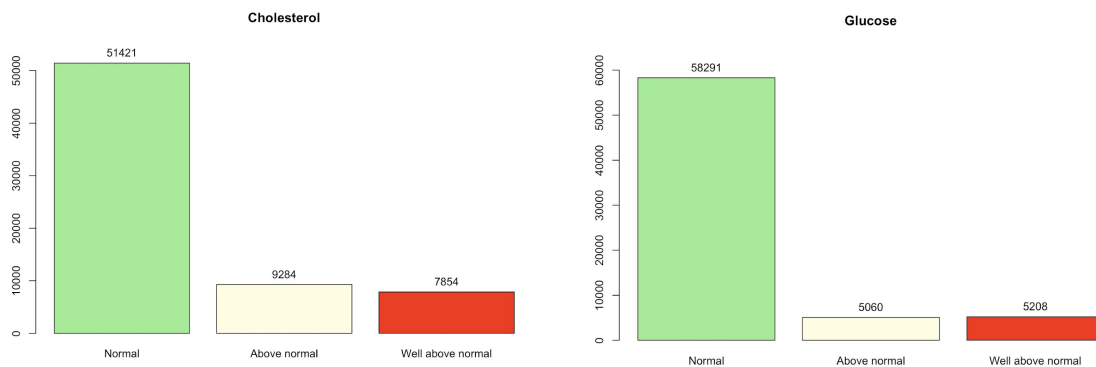
Finally, we examined the correlation between continuous predictors. Below are the correlation plots of the continuous predictors. For the correlation matrix plot: dark blue colors indicate strong positive correlations, dark red is used for strong negative correlations, and white implies no empirical relationship between the predictors.

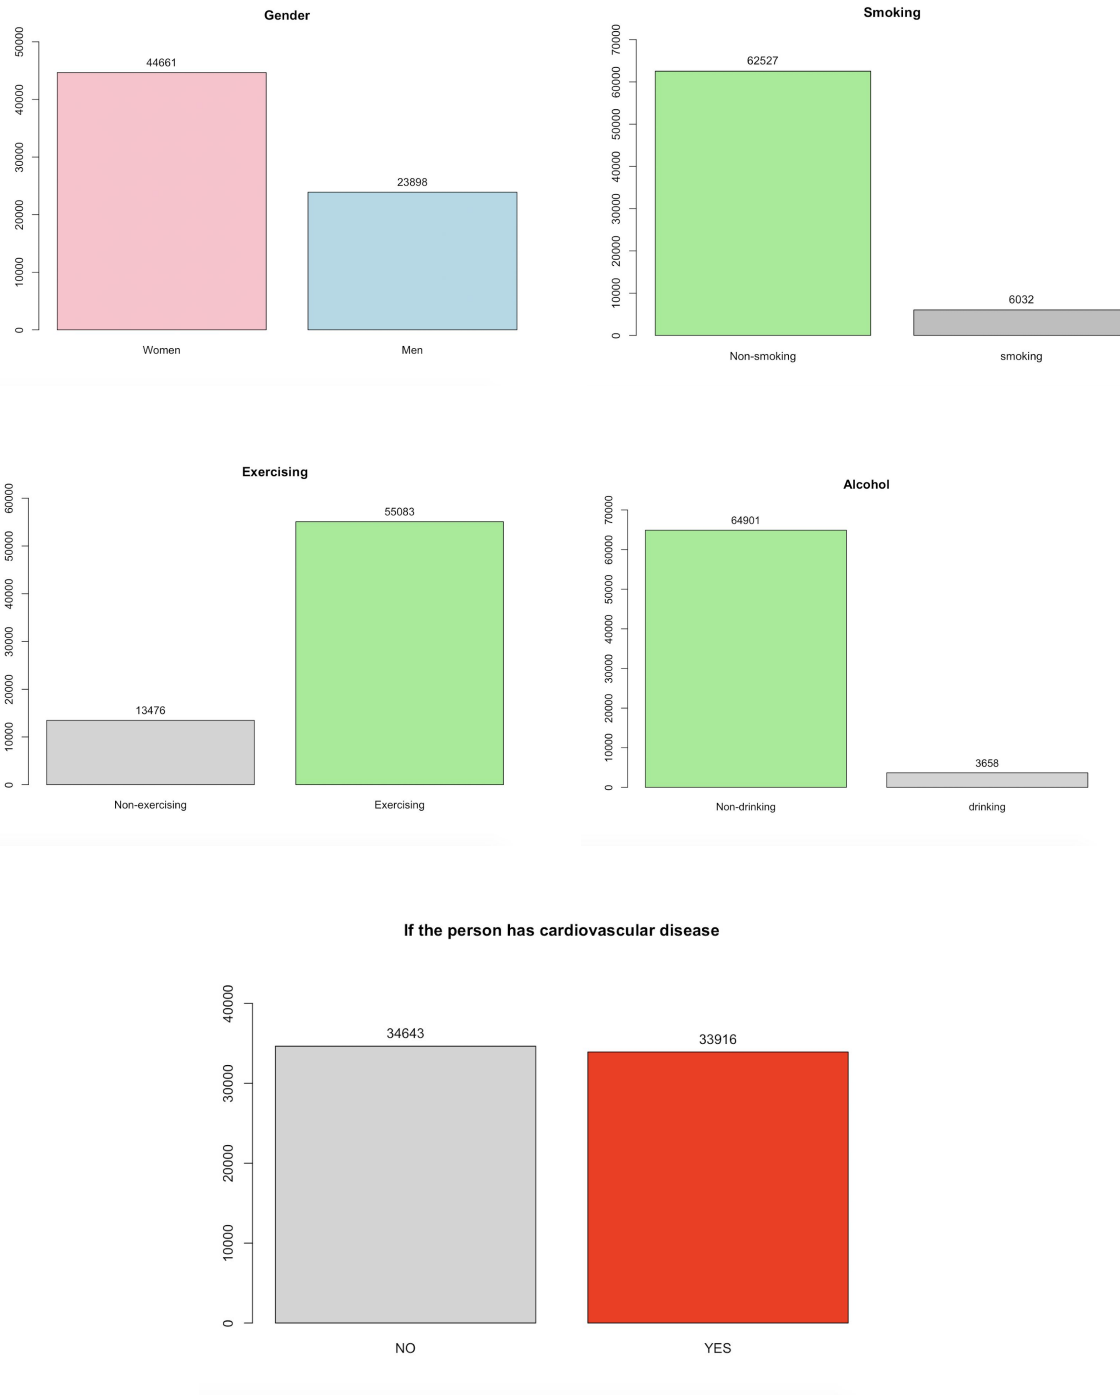


According to the correlation plots displayed above, we see that only “Diastolic” and the transformed “Systolic” variable have a high correlation. This is due to both measures being used to measure blood pressure. To resolve this, we decided to remove “Diastolic” from the analysis. This ensures no multicollinearity and in reality, systolic blood pressure is more important than diastolic blood pressure.

Data Splitting

Prior to model building, we first needed to determine if the dataset was balanced or unbalanced. When examining the bar charts for the categorical variables (Cholesterol, Glucose, Gender, Smoking, Exercising, Alcohol, and Disease) below, we found that the target variable, “Disease,” is pretty balanced.





Based on this, we decided to use the createDataPartition function to split the data into training and testing datasets. The training dataset is often used to build the models while the testing dataset is used solely for validating the performance of the final models. We adopted a simple random sampling technique to split the data into 80% training and 20% in the testing dataset.


Model Building

In this section, we will discuss several models that include both linear and nonlinear classification methods in order to determine the best model in terms of the best predictive ability for the testing data. For this analysis, the following models were built: Logistic Regression, Linear Discriminant Analysis (LDA), Partial Least Squares Discriminant Analysis (PLSDA), Penalized models, Nearest Shrunken Centroids (NSC), Quadratic Discriminant Analysis (QDA), Mixture Discriminant Analysis (MDA), Flexible Discriminant Analysis (FDA), Naive Bayes, K-Nearest Neighbors (KNN), and Neural Network. We were unable to build any Support Vector Machine models for this data due to the following: computation time and computation errors. To determine the best model for predicting CVDs with our predictors, we compared each model using their test ROC curves, the area under these curves, and their test error rates. All models were built using the *train* function from the “caret” package and the same random seed (210).

Using these methods, we first determined the optimal tuning parameters for each model. These parameters are summarized in the table below. If the tuning parameter is marked as NA, this means that the model does not have a tuning parameter.

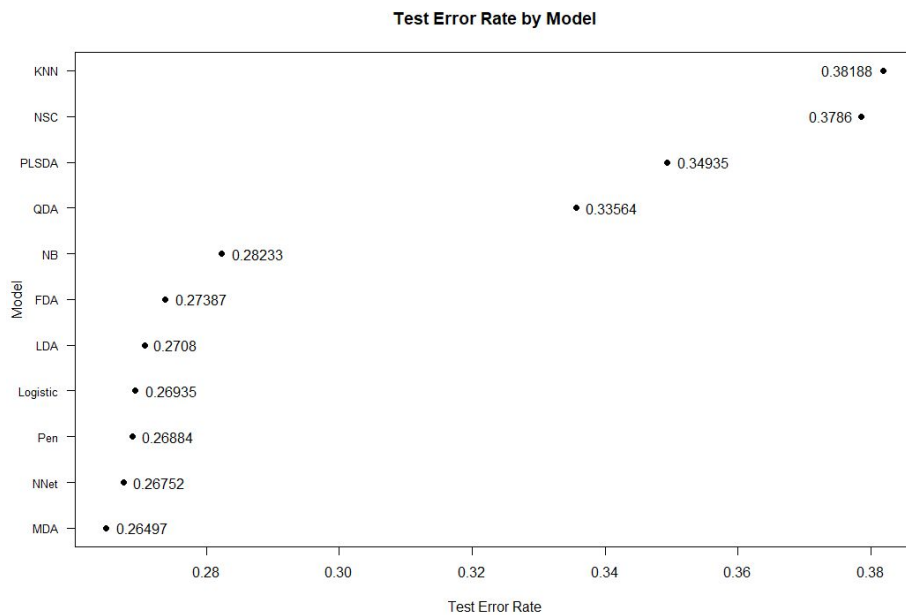
	Tuning Parameters
Logistic	NA
LDA	NA
PLSDA	# Components = 11
Penalized	$\alpha = 0.2, \lambda = 0.01$
NSC	Threshold = 0
QDA	NA
MDA	Subclasses = 2
FDA	Degree = 1, npune = 11
Naïve Bayes	fL = 0, useKernal = True, adjust = 1
KNN	K = 90
NNET	Size = 8, Decay = 0.01

Using these tuning parameters, each model was applied to the test dataset. Confusion matrixes, variable importance, and other summary statistics for each model can be found in Appendix 1. We then compared these eleven models based on their cross-validation statistics which were implemented by using the *resamples* function with models that shared a common set of resampled data sets. The *resamples* function would collect the resampling results of these models into a single object. This could then be used for visualization and/or making some formal comparisons among these models. To further visualize the results, we created the following figure of the cross-validated ROC with a 95% confidence interval across the different models under consideration.

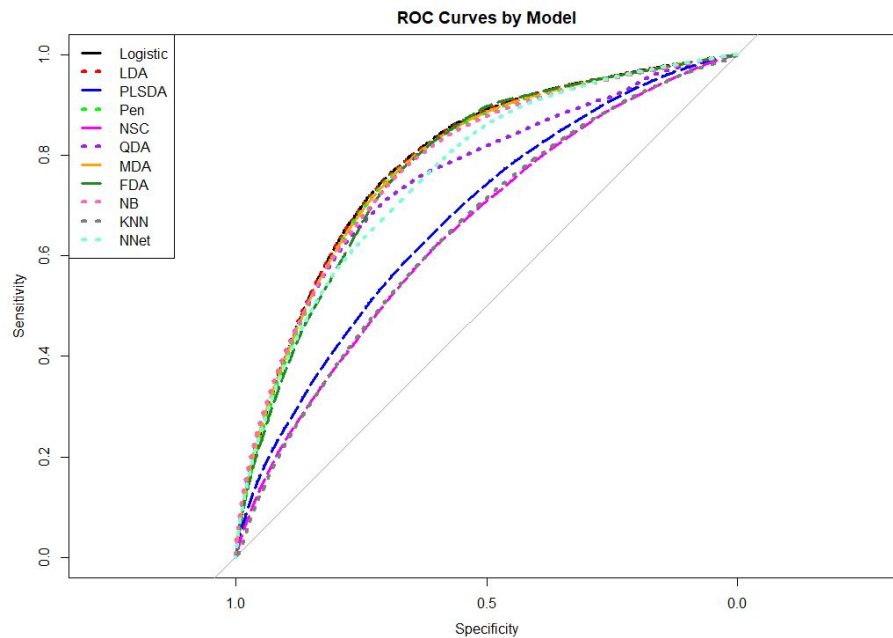


We observe from this figure that the KNN performs the worst and the NNet seems to perform the best in terms of the ROC. In addition, we also observe that other models behave similarly. To further assess possible differences among these 11 models, we compare the test error rates.

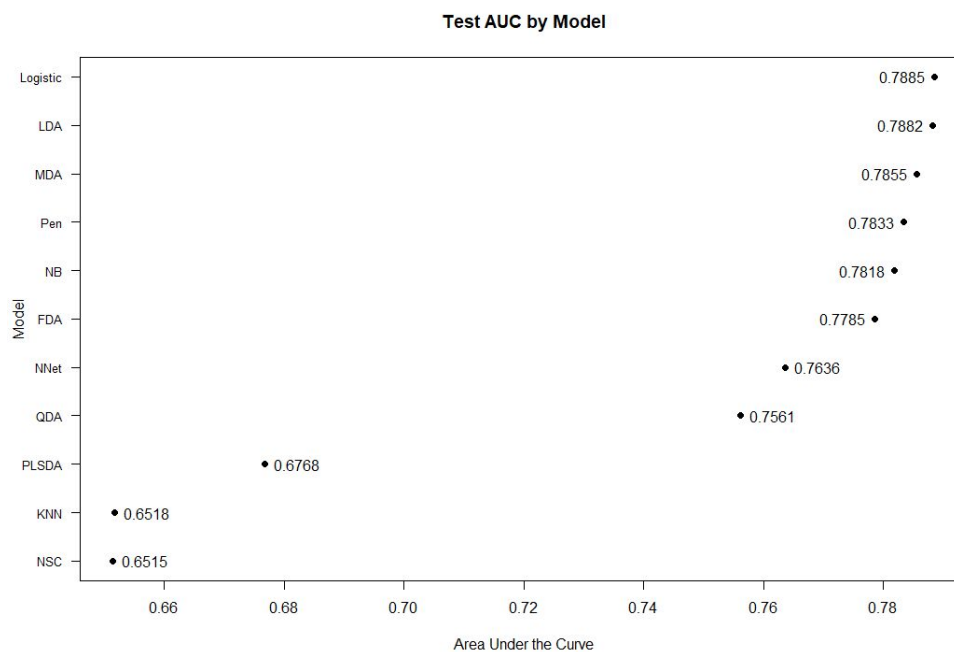
The plot below compares the test error rates for each model. From this plot, we can see that the worst-performing model in regards to the test error rate is KNN and the best is MDA. However, other models are close to MDA, we can see that there is only a less than 1% difference between these models.



The next plot compares the ROC curves between these models. Again, we can see that the models that had a less than 1% difference in the previous plot have very similar ROC curves, with the Logistic Regression model being slightly above the rest.



The final plot below compares the area under the ROC curves for each model. From this plot, we can see the KNN and NSC models have the lowest area while the Logistic Regression model has the highest area, closely followed by other models by a less than 1% difference.



Conclusion

In this project, we analyzed the Cardiovascular Disease dataset provided by kaggle.com, which contains information describing different aspects of patient characteristics, collected at the time of medical examination. Based on our statistical analysis for the data after preprocessing and splitting, we may conclude that the best model would be the Logistic Regression model. While other models may have had slightly better fit statistics, these differences were less than 1%. Therefore, we chose the simpler and more interpretable model as the best choice for this data. Based on Appendix 1, the result of the Logistic Regression model indicates that SystolicTrans and Age are the two most important variables. Also, people who have well-above normal cholesterol level are more likely to have cardiovascular disease.

However, the test error rate for the Logistic Regression model is 26.9%. This error rate may be considered a bit high when being used for medical purposes. Our goal was to build a model with high predictive accuracy, but was unable to do so. This does not mean that this model isn't useful. Our model can be used as a tool by doctors to assist in making a decision on whether or not to run more medical tests to determine whether a patient has a CVD.

We would recommend further investigation for predicting the presence of a cardiovascular disease in a patient. Running further tests could help to improve the high error rate, which in turn would improve the predictability of the models.

Works Cited

- [1] American Heart Association. What is Cardiovascular Disease? 31 May 2017.
<<https://www.heart.org/en/health-topics/consumer-healthcare/what-is-cardiovascular-disease>>.
- [2] Mensah, George A. and David W. Brown. "An Overview Of Cardiovascular Disease Burden In The United States." *Health Affairs* 26.1 (2007): 38-48.
<<https://www.healthaffairs.org/doi/10.1377/hlthaff.26.1.38>>.
- [3] U.S. Department of Health & Human Services. Heart Disease. 2 December 2019.
<<https://www.cdc.gov/heartdisease/facts.htm>>.
- [4] —. Heart Disease and Stroke. 21 March 2019.
<<https://www.cdc.gov/chronicdisease/resources/publications/factsheets/heart-disease-stroke.htm>>.
- [5] World Health Organization. Cardiovascular diseases (CVDs). 17 May 2017.
<[https://www.who.int/news-room/fact-sheets/detail/cardiovascular-diseases-\(cvds\)](https://www.who.int/news-room/fact-sheets/detail/cardiovascular-diseases-(cvds))>.

Appendix 1: Supplementary Material for Models

I. Logistic Regression

```
Generalized Linear Model

54848 samples
 10 predictor
 2 classes: 'No', 'Yes'

No pre-processing
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, 41137, ...
Resampling results:

ROC      Sens      Spec
0.7885347 0.7744169 0.678567

Call:
NULL

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-3.0741  -0.9339  -0.2358   0.9231   3.0430

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  -4.773e+02  5.885e+00 -81.111 < 2e-16 ***
Age           4.935e-02  1.517e-03  32.540 < 2e-16 ***
GenderMale    -1.409e-02  2.480e-02  -0.568 0.569876
Height       -1.531e-02  3.926e-03  -3.899 9.66e-05 ***
WeightTrans   6.670e+00  4.607e-01  14.478 < 2e-16 ***
SystolicTrans 3.773e+02  4.887e+00  77.199 < 2e-16 ***
CholesterolNormal -3.888e-01  3.053e-02 -12.735 < 2e-16 ***
Cholesterolwell Above 7.481e-01  4.721e-02  15.846 < 2e-16 ***
GlucoseNormal -4.713e-02  4.056e-02  -1.162 0.245215
Glucosewell Above -4.483e-01  5.764e-02  -7.778 7.37e-15 ***
SmokeYes     -1.674e-01  3.905e-02  -4.286 1.82e-05 ***
AlcoholYes   -1.831e-01  4.723e-02  -3.876 0.000106 ***
ExerciseYes  -2.254e-01  2.453e-02  -9.186 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)



Null deviance: 76029 on 54847 degrees of freedom
Residual deviance: 61527 on 54835 degrees of freedom
AIC: 61553

Number of Fisher Scoring iterations: 4

logicPred   No  Yes
No    5437 2202
Yes   1491 4581

glm variable importance:

SystolicTrans 100.0000
Age            41.7213
Cholesterolwell Above 19.9365
WeightTrans    18.1511
CholesterolNormal 15.8773
ExerciseYes    11.2457
Glucosewell Above 9.4084
SmokeYes       4.8521
Height         4.3466
AlcoholYes     4.3169
GlucoseNormal  0.7749
GenderMale     0.0000
```



II. Linear Discriminant Analysis

Linear Discriminant Analysis


54848 samples
10 predictor
2 classes: 'No', 'Yes'


No pre-processing
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, 41137, ...
Resampling results:

ROC	Sens	Spec
0.7882096	0.7754388	0.6763379

Tid	Pred	No	Yes
No	5442	2227	
Yes	1486	4556	

	Importance
SystolicTrans	100.0000
Age	53.4291
WeightTrans	40.5428
Cholesterol	10.4023
Exercise	5.0007
Height	2.6500
Smoke	1.4287
Glucose	0.7124
Gender	0.2334
Alcohol	0.0000





III. Partial Least Squares Discriminant Analysis

Partial Least Squares

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, ...
Resampling results across tuning parameters:

ncomp	ROC	Sens	Spec
1	0.6355115	0.5630831	0.6308418
2	0.6487896	0.5954099	0.6150376
3	0.6733142	0.6580427	0.5899808
4	0.6750871	0.6674827	0.5819844
5	0.6757603	0.6646594	0.5864957
6	0.6769890	0.6695670	0.5841899
7	0.6783642	0.6692436	0.5824797
8	0.6810335	0.6729042	0.5849683
9	0.6958609	0.6745958	0.6066696
10	0.6992300	0.6781813	0.6053192
11	0.6998670	0.6779734	0.6068052



ROC was used to select the optimal model using the largest value.
The final value used for the model was ncomp = 11.

```

plsPred  No  Yes
No      4812 2674
Yes     2116 4109

```

	Overall
weightTrans	100.00000
SystolicTrans	28.00463
Cholesterolwell Above	8.38729
CholesterolNormal	7.49798
Glucosewell Above	2.90622
GlucoseNormal	2.38350
SmokeYes	2.09274
AlcoholYes	1.69459
Age	1.32596
ExerciseYes	1.11429
Height	0.06204
GenderMale	0.00000



IV. Penalized Model

glmnet

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, ...

Resampling results across tuning parameters:

alpha	lambda	ROC	Sens	Spec
0.0	0.01000000	0.7882030	0.7726039	0.6790624
0.0	0.02357143	0.7881643	0.7724654	0.6790034
0.0	0.03714286	0.7878861	0.7720843	0.6791626
0.0	0.05071429	0.7875939	0.7714203	0.6788972
0.0	0.06428571	0.7872977	0.7711028	0.6784903
0.0	0.07785714	0.7870018	0.7709700	0.6778888
0.0	0.09142857	0.7867104	0.7707333	0.6771340
0.0	0.10500000	0.7864231	0.7706236	0.6760902
0.0	0.11857143	0.7861461	0.7706755	0.6753295
0.0	0.13214286	0.7858769	0.7707737	0.6743978
0.0	0.14571429	0.7856144	0.7708314	0.6736017
0.0	0.15928571	0.7853601	0.7708776	0.6728645
0.0	0.17285714	0.7851133	0.7711432	0.6720035
0.0	0.18642857	0.7848717	0.7713799	0.6713254
0.0	0.20000000	0.7846380	0.7717206	0.6703524
0.1	0.01000000	0.7884570	0.7737471	0.6783075
0.1	0.02357143	0.7882553	0.7731582	0.6772873
0.1	0.03714286	0.7879644	0.7728522	0.6764676
0.1	0.05071429	0.7876048	0.7727598	0.6754415
0.1	0.06428571	0.7871934	0.7727945	0.6745924
0.1	0.07785714	0.7868470	0.7729619	0.6739142
0.1	0.09142857	0.7865387	0.7730600	0.6730591
0.1	0.10500000	0.7862442	0.7732621	0.6723810
0.1	0.11857143	0.7859635	0.7732044	0.6719092
0.1	0.13214286	0.7857468	0.7735508	0.6711662
0.1	0.14571429	0.7855760	0.7740820	0.6702875



0.1	0.15928571	0.7854317	0.7745497	0.6691552
0.1	0.17285714	0.7853318	0.7750058	0.6682412
0.1	0.18642857	0.7852818	0.7755774	0.6674864
0.1	0.20000000	0.7852642	0.7762125	0.6667905
0.2	0.01000000	0.7884757	0.7741455	0.6776824
0.2	0.02357143	0.7881513	0.7738915	0.6759074
0.2	0.03714286	0.7875656	0.7743187	0.6739791
0.2	0.05071429	0.7870476	0.7748499	0.6728173
0.2	0.06428571	0.7866450	0.7753176	0.6716379
0.2	0.07785714	0.7863617	0.7760046	0.6699455
0.2	0.09142857	0.7862180	0.7767148	0.6686304
0.2	0.10500000	0.7862071	0.7777309	0.6674628
0.2	0.11857143	0.7862318	0.7785508	0.6664662
0.2	0.13214286	0.7862605	0.7795958	0.6653634
0.2	0.14571429	0.7862889	0.7805196	0.6643196
0.2	0.15928571	0.7863119	0.7819053	0.6635058
0.2	0.17285714	0.7863283	0.7832737	0.6624738
0.2	0.18642857	0.7863302	0.7844284	0.6613593
0.2	0.20000000	0.7863141	0.7854965	0.6601504
0.4	0.01000000	0.7883970	0.7747864	0.6760725
0.4	0.02357143	0.7873840	0.7758891	0.6725343
0.4	0.03714286	0.7866489	0.7775058	0.6696447
0.4	0.05071429	0.7863399	0.7792841	0.6670441
0.4	0.06428571	0.7862272	0.7812875	0.6645968
0.4	0.07785714	0.7860328	0.7836836	0.6612944
0.4	0.09142857	0.7857126	0.7863799	0.6572018
0.4	0.10500000	0.7852068	0.7889607	0.6531800
0.4	0.11857143	0.7844617	0.7916744	0.6495651
0.4	0.13214286	0.7834096	0.7949018	0.6460150
0.4	0.14571429	0.7823942	0.7965993	0.6440277
0.4	0.15928571	0.7813212	0.8015185	0.6381837
0.4	0.17285714	0.7801439	0.8056351	0.6299455
0.4	0.18642857	0.7788565	0.8091109	0.6196609
0.4	0.20000000	0.7774928	0.8067206	0.6197081

ROC was used to select the optimal model using the largest value.

The final values used for the model were alpha = 0.2 and lambda = 0.01.

```
penPpred  No  Yes
No  5445 2200
Yes 1480 4680
```

	Overall
SystolicTrans	1.000e+02
WeightTrans	1.787e+00
CholesterolWell Above	1.816e-01
CholesterolNormal	1.111e-01
GlucoseWell Above	7.717e-02
ExerciseYes	5.382e-02
SmokeYes	3.993e-02
AlcoholYes	3.870e-02
Age	1.337e-02
Height	3.560e-03
GenderMale	5.445e-04
GlucoseNormal	0.000e+00



V. Nearest Shrunk Centroids

Nearest Shrunk Centroids

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, 41137, ...



Resampling results across tuning parameters:

threshold	ROC	Sens	Spec
0	0.6659480	0.6240127	0.6136282
1	0.6655163	0.6186085	0.6178682
2	0.6649961	0.6131062	0.6222969
3	0.6642235	0.6090647	0.6255403
4	0.6633550	0.6068072	0.6267197
5	0.6628701	0.6071882	0.6260062
6	0.6621705	0.6087182	0.6243255
7	0.6614166	0.6095381	0.6225800
8	0.6606440	0.6104965	0.6203745
9	0.6595232	0.6110450	0.6189828
10	0.6582969	0.6115878	0.6177503
11	0.6573506	0.6122229	0.6168362
12	0.6566522	0.6134238	0.6137815
13	0.6557926	0.6148037	0.6098423
14	0.6542360	0.6156178	0.6055315
15	0.6522387	0.6158256	0.6035029
16	0.6506552	0.6151790	0.6033908
17	0.6491252	0.6142436	0.6012207
18	0.6463103	0.6122806	0.5986024
19	0.6432750	0.6078233	0.5994575
20	0.6421701	0.6035508	0.6025004
21	0.6412549	0.5992321	0.6052012
22	0.6397251	0.5966975	0.6046410
23	0.6371071	0.5953580	0.6015627
24	0.6348986	0.5963395	0.5970338
25	0.6348838	0.5991051	0.5937314

ROC was used to select the optimal model using the largest value.
The final value used for the model was threshold = 0.

nsCpred	No	Yes
No	4386	2649
Yes	2542	4134

	Importance
Age	97.869807
CholesterolNormal	48.316269
CholesterolWell Above	38.712981
GlucoseNormal	18.393890
GlucoseWell Above	11.502461
ExerciseYes	7.089080
Height	4.169910
WeightTrans	3.789643
SmokeYes	2.523441
AlcoholYes	0.213218
GenderMale	0.005682
SystolicTrans	0.000000



VI. Quadratic Discriminant Analysis

Quadratic Discriminant Analysis



54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, 41137, ...
Resampling results:

ROC	Sens	Spec
0.7562116	0.8057217	0.5288751

qdaPred	No	Yes
No	5589	3260
Yes	1339	3520

	Importance
SystolicTrans	100.0000
Age	53.4291
WeightTrans	40.5428
Cholesterol	10.4023
Exercise	5.0007
Height	2.6500
Smoke	1.4287
Glucose	0.7124
Gender	0.2334
Alcohol	0.0000



VII. Mixture Discriminant Analysis

Mixture Discriminant Analysis

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, ...


Resampling results across tuning parameters:


subclasses	ROC	Sens	Spec
1	0.7882096	0.7754388	0.6763497
2	0.7902966	0.7661547	0.6888103
3	0.7902874	0.7654388	0.6885213
4	0.7884634	0.7671363	0.6851953
5	0.7855521	0.7618129	0.6839805
6	0.7859941	0.7639145	0.6848533
7	0.7830827	0.7601501	0.6822173
8	0.7833445	0.7566224	0.6857143
9	0.7810281	0.7465185	0.6907209
10	0.7797384	0.7430774	0.6906089

ROC was used to select the optimal model using the largest value.
The final value used for the model was subclasses = 2.

mdaPred	No	Yes
No	5412	2117
Yes	1516	4666

	Importance
SystolicTrans	100.0000
Age	53.4291
WeightTrans	40.5428
Cholesterol	10.4023
Exercise	5.0007
Height	2.6500
Smoke	1.4287
Glucose	0.7124
Gender	0.2334
Alcohol	0.0000





VIII. Flexible Discriminant Analysis

Flexible Discriminant Analysis

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing



Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, 41137, ...
Resampling results across tuning parameters:

nprune	ROC	Sens	Spec
2	0.7500078	0.8047864	0.6202211
6	0.7860634	0.7981755	0.6425947
11	0.7927760	0.8060450	0.6406664

Tuning parameter 'degree' was held constant at a value of 1.
ROC was used to select the optimal model using the largest value.
The final values used for the model were degree = 1 and nprune = 11.

	No	Yes
fdaPred		
No	5621	2448
Yes	1307	4335

	Overall
SystolicTrans	100.000
Age	46.280
'CholesterolWellAbove'	27.641
CholesterolNormal	19.262
WeightTrans	14.408
ExerciseYes	8.858
GenderMale	0.000
'GlucoseWellAbove'	0.000
SmokeYes	0.000
GlucoseNormal	0.000
Height	0.000
AlcoholYes	0.000



IX. Naive Bayes

Naive Bayes

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, ...

Resampling results across tuning parameters:

usekernel	ROC	Sens	Spec
FALSE	0.7824492	0.7672748	0.6674805
TRUE	0.7868560	0.8285855	0.6033731

Tuning parameter 'fL' was held constant at a value of 0

Tuning parameter


'adjust' was held constant at a value of 1


ROC was used to select the optimal model using the largest value.

The final values used for the model were fL = 0, usekernel = TRUE and adjust = 1.

nbPried	No	Yes
No	5771	2714
Yes	1157	4069

	Importance
SystolicTrans	100.0000
Age	53.4291
WeightTrans	40.5428
Cholesterol	10.4023
Exercise	5.0007
Height	2.6500
Smoke	1.4287
Glucose	0.7124
Gender	0.2334
Alcohol	0.0000





X. K-Nearest Neighbors

k-Nearest Neighbors

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, ...

Resampling results across tuning parameters:



k	ROC	Sens	Spec
1	0.5531709	0.5581178	0.5482353
2	0.5766450	0.5616282	0.5444612
3	0.5915298	0.5819111	0.5608374
4	0.6022008	0.5823961	0.5608905
5	0.6107851	0.5971420	0.5681144
6	0.6173465	0.5973961	0.5673655
7	0.6230811	0.6067263	0.5720655
8	0.6274421	0.6070439	0.5710512
9	0.6308207	0.6146132	0.5755919
10	0.6337767	0.6142379	0.5756627
11	0.6366822	0.6199076	0.5779095
12	0.6390669	0.6197402	0.5770662
13	0.6407544	0.6237760	0.5783518
14	0.6422975	0.6238106	0.5763467
15	0.6437673	0.6282737	0.5762229
16	0.6450278	0.6290069	0.5761286
17	0.6465473	0.6326501	0.5754681
18	0.6478803	0.6338222	0.5745658
19	0.6490043	0.6378522	0.5746307
20	0.6499920	0.6391339	0.5737343
21	0.6509069	0.6409296	0.5735515
22	0.6516753	0.6424423	0.5721480
23	0.6524167	0.6454330	0.5729146
24	0.6529511	0.6451328	0.5722070
25	0.6534494	0.6465012	0.5712811
26	0.6541318	0.6475289	0.5709627
27	0.6547561	0.6501039	0.5708625
28	0.6551820	0.6503637	0.5702727

29	0.6555215	0.6524018	0.5693823
30	0.6561211	0.6529619	0.5692113
31	0.6568358	0.6545958	0.5685154
32	0.6573352	0.6546247	0.5675247
33	0.6577442	0.6563395	0.5673714
34	0.6581379	0.6560450	0.5676544
35	0.6585870	0.6577829	0.5666460
36	0.6588742	0.6577367	0.5662273
37	0.6591750	0.6594053	0.5664868
38	0.6595728	0.6595035	0.5665635
39	0.6598742	0.6611721	0.5656140
40	0.6601957	0.6617783	0.5648651
41	0.6605810	0.6626674	0.5652779
42	0.6607394	0.6633256	0.5647943
43	0.6609734	0.6646074	0.5644287
44	0.6613057	0.6646074	0.5636031
45	0.6615033	0.6658314	0.5629368
46	0.6618434	0.6660912	0.5630134
47	0.6619942	0.6674365	0.5625534
48	0.6620865	0.6673383	0.5623942
49	0.6621940	0.6691166	0.5619991
50	0.6623479	0.6685508	0.5615922
51	0.6624864	0.6699423	0.5614035
52	0.6626081	0.6700693	0.5611381
53	0.6627521	0.6709122	0.5609789
54	0.6628362	0.6715069	0.5601297
55	0.6629219	0.6724192	0.5598939
56	0.6630207	0.6721536	0.5595518
57	0.6631615	0.6731351	0.5599351
58	0.6632705	0.6735046	0.5595223
59	0.6633263	0.6738857	0.5589149
60	0.6634235	0.6744977	0.5578888
61	0.6635179	0.6751559	0.5581660
62	0.6637002	0.6756871	0.5573994
63	0.6638416	0.6764896	0.5570809
64	0.6639803	0.6769053	0.5574230
65	0.6641498	0.6779677	0.5566033
66	0.6642961	0.6785855	0.5563438
67	0.6643954	0.6788453	0.5559959
68	0.6644556	0.6788972	0.5555713
69	0.6644882	0.6796536	0.5557718
70	0.6644764	0.6800635	0.5552410
71	0.6644738	0.6804273	0.5543388
72	0.6644292	0.6803522	0.5536370
73	0.6644647	0.6811085	0.5536960
74	0.6644741	0.6812125	0.5528291
75	0.6644047	0.6816628	0.5521333
76	0.6644047	0.6815878	0.5515436
77	0.6643363	0.6826443	0.5505823
78	0.6643100	0.6829965	0.5504231
79	0.6643266	0.6835739	0.5498039
80	0.6643931	0.6842725	0.5497567
81	0.6644572	0.6845381	0.5490314
82	0.6645029	0.6853811	0.5489842
83	0.6645222	0.6856120	0.5481350
84	0.6645514	0.6856640	0.5476633
85	0.6645801	0.6862298	0.5470441
86	0.6646461	0.6867206	0.5467669
87	0.6646042	0.6871247	0.5456347
88	0.6646831	0.6872460	0.5458824
89	0.6646797	0.6875520	0.5456995
90	0.6647307	0.6874827	0.5455285
91	0.6647263	0.6881524	0.5446617
92	0.6647002	0.6885508	0.5446970
93	0.6647141	0.6883891	0.5448327
94	0.6646666	0.6887009	0.5439304
95	0.6646077	0.6891686	0.5437063
96	0.6645802	0.6895266	0.5436356
97	0.6645058	0.6895670	0.5436120
98	0.6645113	0.6896998	0.5432051
99	0.6645242	0.6903522	0.5427746
100	0.6644779	0.6901617	0.5420610

ROC was used to select the optimal model using the largest value.
The final value used for the model was k = 90.

konfPred	No	Yes
No	4732	3040
Yes	2196	3743

	Importance
SystolicTrans	100.0000
Age	53.4291
WeightTrans	40.5428
Cholesterol	10.4023
Exercise	5.0007
Height	2.6500
Smoke	1.4287
Glucose	0.7124
Gender	0.2334
Alcohol	0.0000



XI. Neural Network

Neural Network

54848 samples
10 predictor
2 classes: 'No', 'Yes'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 41137, 41137, 41137, 41137, 41137, 41137, ...



Resampling results across tuning parameters:

size	decay	ROC	Sens	Spec
1	0.00	0.5643346	0.9382564	0.1676367
1	0.01	0.6647351	0.8248961	0.4378358
1	0.10	0.7570364	0.7750635	0.6238420
1	0.50	0.7102799	0.7933545	0.5429102
2	0.00	0.5554200	0.9263164	0.1687336
2	0.01	0.7501860	0.7915069	0.5859590
2	0.10	0.7371549	0.7891975	0.5795312
2	0.50	0.7788922	0.7379619	0.6884682
3	0.00	0.5821396	0.9155831	0.2233820
3	0.01	0.7679598	0.7846478	0.6151732
3	0.10	0.7824676	0.7668476	0.6683238
3	0.50	0.7630939	0.7475404	0.6630871
4	0.00	0.6087686	0.9015820	0.2761787
4	0.01	0.7720788	0.7694400	0.6484269
4	0.10	0.7906050	0.7537644	0.7002565
4	0.50	0.7729096	0.7494804	0.6631107
5	0.00	0.6414549	0.8745092	0.3558131
5	0.01	0.7937158	0.7477425	0.7104349
5	0.10	0.7910820	0.7543187	0.7000737
5	0.50	0.7811261	0.7402252	0.6908801
6	0.00	0.6968653	0.8227945	0.4914168
6	0.01	0.7844264	0.7580312	0.6818694
6	0.10	0.7913389	0.7538626	0.7008816
6	0.50	0.7805268	0.7387413	0.6911632
7	0.00	0.7099556	0.8166570	0.5227893
7	0.01	0.7833756	0.7534700	0.6849005
7	0.10	0.7914762	0.7524654	0.7020257
7	0.50	0.7743754	0.7497113	0.6642430
8	0.00	0.7436450	0.8144284	0.5661448
8	0.01	0.7943617	0.7462413	0.7127466
8	0.10	0.7821807	0.7626386	0.6736488
8	0.50	0.7811289	0.7400000	0.6914522
9	0.00	0.7383971	0.8021189	0.5738110
9	0.01	0.7727820	0.7669630	0.6551732
9	0.10	0.7909299	0.7528811	0.7011706
9	0.50	0.7811614	0.7404273	0.6910865
10	0.00	0.7465481	0.7814723	0.6114699
10	0.01	0.7941097	0.7460566	0.7129707
10	0.10	0.7909446	0.7520843	0.7017544
10	0.50	0.7802166	0.7382333	0.6915465

ROC was used to select the optimal model using the largest value.
The final values used for the model were size = 8 and decay = 0.01.

```
nnnetPred  No  Yes
No  5303 2043
Yes 1625 4740
```

	Overall
SystolicTrans	100.000
Height	37.546
Age	36.327
Cholesterolwell Above	32.433
Glucosewell Above	21.317
WeightTrans	17.670
GenderMale	10.784
CholesterolNormal	7.579
GlucoseNormal	6.972
SmokeYes	1.917
AlcoholYes	0.183
ExerciseYes	0.000



R Code

Data Algorithms II

Project Code

Data Cleaning

Import data

```
data <- read.csv("cardio.csv")
```

```
View(data)
```

The purpose of this code is to clean the data and apply unit conversions from metri to imperial

ID column can be removed

```
data <- data[,-1]
```

Create better column names

```
varNames <- c("Age", "Gender", "Height", "Weight", "Systolic", "Diastolic",  
             "Cholesterol", "Glucose", "Smoke", "Alcohol", "Exercise", "Disease")
```

```
colnames(data) <- varNames
```

Age is in days, Height is in cm, Weight is in kg.

Change these to years, inches, and pounds, respectively.

```
data$Age <- round(data$Age / 365.25, 2) ### This accounts for leap years
```

```
data$Height <- round(data$Height / 2.54, 2)
```

```
data$Weight <- round(data$Weight * 2.20462, 2)
```

Gender: 1 = women, 2 = men.

Cholesterol: 1 = normal, 2 = above normal, 3 = well above normal

Glucose: 1 = normal, 2 = above normal, 3 = well above normal

Subtract 1 from everyone of these variables to have a base level of zero

```
data$Gender <- data$Gender - 1
```

```
data$Cholesterol <- data$Cholesterol - 1
```

```
data$Glucose <- data$Glucose - 1
```

Now convert all categorical variables to factor type

```
data[,c(2,7:12)] <- lapply(data[,c(2,7:12)], factor)
```

Output new csv

```
write.csv(data, "cardio2.csv")
```

```
rm(list = ls())
```

```
data <- read.csv("cardio2.csv")
data <- data[,-1]
data[,c(2,7:12)] <- lapply(data[,c(2,7:12)], factor)
```

```
#### Grab continuous variables
contin <- data[,c(1,3,4,5,6)]
```

```
#### Look at boxplots
```

```
par(mfrow=c(1,5))
for (i in 1:5){
  boxplot(contin[,i], xlab = colnames(contin)[i])
}
```

```
#### We can see that systolic and diastolic have massive outliers that make no physical sense.
Why is this?
```

```
dev.off()
plot(contin$Systolic, contin$Diastolic)
```

```
#### There are plenty of assumptions we could make to fix the data to our needs
#### However, there are just too many that are needed for our analysis to still be valid
#### Instead, we will focus on the BP range from hypotension to hypertension
```

```
data2 <- data[which(data$Systolic >= 50 & data$Systolic < 220),]
data2 <- data2[which(data2$Diastolic >= 20 & data2$Diastolic < 190),]
```

```
plot(data2$Systolic, data2$Diastolic)
```

```
#### Remove cases where Diastolic is greater than systolic
```

```
data2 <- data2[-which(data2$Diastolic > data2$Systolic),]
plot(data2$Systolic, data2$Diastolic)
```

```
#### Now lets look at height
quantile(data2$Height)
```

```
#### Someone who suffers from Dwarfism has a height below 4ft 10in.
```

```
### For this study, we will look at subject's who are greater than 4ft
data2 <- data2[which(data2$Height >= 48),]
```

```
### Weight has some illogical values for adults.
```

```
### Use values of weight above 80lb
```

```
data2 <- data2[which(data2$Weight >= 80),]
```

```
### Look at boxplots again
```

```
par(mfrow=c(1,5))
```

```
for (i in 1:5){
```

```
  boxplot(contain2[,i], xlab = colnames(contain2)[i])
```

```
}
```

```
### How much was removed?
```

```
1 - nrow(data2)/nrow(data)
```

```
### Only 2.05%. This is good.
```

```
### Output new csv
```

```
write.csv(data2, "cardio3.csv")
```

```
rm(list = ls())
```

```
##### Data Preprocessing/ Exploration #####
```

```
### Import Data
```

```
data <- read.csv("cardio3.csv")
```

```
data <- data[,-1]
```

```
data[,c(2,7:12)] <- lapply(data[,c(2,7:12)], factor)
```

```
library(caret)
```

```
library(corrplot)
```

```
library(e1071)
```

```
##### Continuous Variable #####
```

```
### Scatter plots
```

```

contin <- data[,c(1,3,4,5,6)]
plot(contin, col = as.numeric(data$Disease)+1,
     cex = 0.5, pch = 16)

#### Corr Plots
par(mfrow=c(1,2))
corplot(cor(contin))
corplot(cor(contin), method = "number")

#### Histograms
par(mfrow = c(2,3))
for(i in 1:5){
  hist(contin[,i], xlab = colnames(contin)[i], main = colnames(contin)[i])
}

#### Box Plots
par(mfrow=c(1,5))
for(i in 1:5){
  boxplot(contin[,i], xlab = colnames(contin)[i])
}

#### Skewness

apply(contin, 2, skewness)
for(i in 1:5){
  print(abs(max(contin[i])/min(contin[i])))
}

#### Possible transformations

apply(contin, 2, BoxCoxTrans)

#### Quadratic transformation for Age
#### Square-root transformation for Height
#### Inverse square-root transformation for Weight
#### Inverse transformation for Systolic
#### No transformation for Diastolic

continTrans <- data.frame(contin$Age ^ 2, contin$Height ^ 0.5,
                          contin$Weight ^ -0.5, contin$Systolic ^ -1,
                          contin$Diastolic)

dev.off()

```

```

plot(continTrans, col = as.numeric(data$Disease)+1,
     cex = 0.5, pch = 16)

par(mfrow = c(2,3))
for(i in 1:5){
  hist(continTrans[,i], xlab = colnames(continTrans)[i])
}
apply(continTrans, 2, skewness)

corrplot(cor(continTrans))
corrplot(cor(continTrans), method = "number")

### Not much changed.
### Try center and scaling

continPP <- preProcess(contin, method = c("scale","center"))
contin2 <- predict(continPP, contin)

dev.off()
plot(contin2, col = as.numeric(data$Disease)+1,
     cex = 0.5, pch = 16)

### Didn't change much.
### Try BoxCox, center and scaling

continPP <- preProcess(contin, method = c("scale","center","BoxCox"))
contin3 <- predict(continPP, contin)

plot(contin3, col = as.numeric(data$Disease)+1,
     cex = 0.5, pch = 16)

### Has increased the separation of the target variable

par(mfrow = c(2,3))
for(i in 1:5){
  hist(contin3[,i], xlab = colnames(contin3)[i])
}
apply(contin3, 2, skewness)

### This looks good now

par(mfrow=c(1,2))
corrplot(cor(contin3))

```

```
corrplot(cor(contin3), method = "number")
```

```
#### Correlations haven't really change.
```

```
#### How much can one pair of correlated variable mess with classification?
```

```
#### Use BoxCox only
```

```
##### Categorical Variable #####
```

```
##Gender
```

```
Bgender <- barplot((table(data$Gender)),  
  names.arg=c("Women", "Men"),  
  col = c("Pink", "lightblue"),  
  ylim = c(0, 50000),  
  main = "Gender")
```

```
text(x=Bgender, y= table(data$Gender),  
  labels=as.character(table(data$Gender)),  
  pos = 3,  
  col = "Black")
```

```
##Cholesterol
```

```
BCholesterol <- barplot((table(data$Cholesterol)),  
  names.arg=c("Normal", "Above normal", "Well above normal"),  
  col = c("lightgreen", "lightyellow", "red"),  
  ylim = c(0, 55000),  
  main = "Cholesterol")
```

```
text(x=BCholesterol, y= table(data$Cholesterol),  
  labels=as.character(table(data$Cholesterol)),  
  pos = 3,  
  col = "Black")
```

```
##Glucose
```

```
BGlucose <- barplot((table(data$Glucose)),  
  names.arg=c("Normal", "Above normal", "Well above normal"),  
  col = c("lightgreen", "lightyellow", "red"),  
  ylim = c(0, 65000),  
  main = "Glucose")
```

```
text(x=BGlucose, y= table(data$Glucose),  
  labels=as.character(table(data$Glucose)),  
  pos = 3,  
  col = "Black")
```

```

##Smoke
BSmoke <- barplot((table(data$Smoke)),
  names.arg=c("Non-smoking", "smoking"),
  col = c("lightgreen", "gray"),
  ylim = c(0, 70000),
  main = "Smoking")
text(x=BSmoke, y= table(data$Smoke),
  labels=as.character(table(data$Smoke)),
  pos = 3,
  col = "Black")

##Alcohol
BAlcohol <- barplot((table(data$Alcohol)),
  names.arg=c("Non-drinking", "drinking"),
  col = c("lightgreen", "lightgray"),
  ylim = c(0, 70000),
  main = "Alcohol")
text(x=BAlcohol , y= table(data$Alcohol ),
  labels=as.character(table(data$Alcohol)),
  pos = 3,
  col = "Black")

##Exercise
BExercise <- barplot((table(data$Exercise)),
  names.arg=c("Non-exercising", "Exercising"),
  col = c("lightgray", "lightgreen"),
  ylim = c(0, 60000),
  main = "Exercising")
text(x=BExercise , y= table(data$Exercise),
  labels=as.character(table(data$Exercise)),
  pos = 3,
  col = "Black")

##Disease
BDisease <- barplot((table(data$Disease)),
  names.arg=c("NO", "YES"),
  col = c("lightgray", "red"),
  ylim = c(0, 45000),
  main = "If the person has cardiovascular disease")
text(x=BDisease , y= table(data$Disease),
  labels=as.character(table(data$Disease)),

```

```

    pos = 3,
    col = "Black")

rm(list = ls())

##### MODELING #####

### Import Data
data <- read.csv("cardio3.csv")
data <- data[,-1]
data[,c(2,7:12)] <- lapply(data[,c(2,7:12)], factor)
#View(data)
#str(data)

library(caret)
library(corrplot)
library(e1071)

### Data Pre-Processing

### Remove Diastolic
data <- data[,-6]

par(mfrow=c(2,2))
hist(data$Weight, xlab = "Weight", main = "Weight Original")

### Apply transformations to Weight and Systolic
weightPP <- BoxCoxTrans(data$Weight)
WeightTrans <- predict(weightPP, data$Weight)

hist(WeightTrans, xlab = "Weight Tans", main = "Weight Transformed")

hist(data$Systolic, xlab = "Systolic", main = "Systolic Original")

systolicPP <- BoxCoxTrans(data$Systolic)
SystolicTrans <- predict(systolicPP, data$Systolic)

hist(SystolicTrans, xlab = "Systolic Trans", main = "Systolic Transformed")

### Put transformations in data

```



```
data$Weight <- WeightTrans
data$Systolic <- SystolicTrans
```

```
### Update column names
colnames(data)[c(4,5)] <- c("WeightTrans", "SystolicTrans")
```

```
par(mfrow=c(1,2))
corrplot(cor(data[,c(1,3:6)]))
corrplot(cor(data[,c(1,3:6)]), method = "number")
```

```
### Make a copy of the data.
### One will hold numeric factors
### One will hold string factors
```

```
data2 <- data
```

```
### Make categorical variables strings
data2$Gender <- ifelse(data$Gender == 0, "Female", "Male")
data2$Alcohol <- ifelse(data$Alcohol == 0, "No", "Yes")
data2$Smoke <- ifelse(data$Smoke == 0, "No", "Yes")
data2$Exercise <- ifelse(data$Exercise == 0, "No", "Yes")
data2$Disease <- ifelse(data$Disease == 0, "No", "Yes")
```

```
data2$Cholesterol <- ifelse(data$Cholesterol == 0, "Normal",
                           ifelse(data$Cholesterol == 1, "Above", "Well Above"))
data2$Glucose <- ifelse(data$Glucose == 0, "Normal",
                       ifelse(data$Glucose == 1, "Above", "Well Above"))
```

```
data2[,c(2,6:11)] <- lapply(data2[,c(2,6:11)], factor)
```

```
### Create control function
set.seed(210)
ctrl <- trainControl(method = "LGOCV",
                     summaryFunction = twoClassSummary,
                     classProbs = TRUE,
                     savePredictions = TRUE)
```

```
### Create training and testing data
```

```

set.seed(210)
inTrain <- createDataPartition(data$Disease, p = 0.8)[[1]]
Train <- data[inTrain,]
Test <- data[-inTrain,]

Train2 <- data2[inTrain,]
Test2 <- data2[-inTrain,]

#### Logistic Regression
set.seed(210)
logicTune <- train(x = Train2[,c(1:10)], y = Train2$Disease,
  method = "glm",
  metric = "ROC",
  trControl = ctrl)
logicTune

summary(logicTune)

logicPred <- predict(logicTune, Test2)
table(logicPred, Test2$Disease)

#### Test error rate
mean(logicPred != Test2$Disease)

#### Importance
varImp(logicTune)
plot(varImp(logicTune))

#### Test ROC
library(pROC)

logicROC <- roc(response = logicTune$pred$obs,
  predictor = logicTune$pred$Yes,
  levels = rev(levels(logicTune$pred$obs)))
plot(logicROC, legacy.axes = TRUE)
auc(logicROC)

#### Save results

Test_Error <- c(mean(logicPred != Test2$Disease))
Test_AUC <- c(0.7885)
Models <- c("Logistic")

```

```

#### LDA
set.seed(210)
ldaTune <- train(form = Disease~., data = Train2,
                 method = "lda",
                 metric = "ROC",
                 trControl = ctrl)
ldaTune

ldaPred <- predict(ldaTune, Test2)
table(ldaPred, Test2$Disease)

#### Test error rate
mean(ldaPred != Test2$Disease)

#### Importance
varImp(ldaTune)
plot(varImp(ldaTune))

#### Test ROC
library(pROC)

ldaROC <- roc(response = ldaTune$pred$obs,
              predictor = ldaTune$pred$Yes,
              levels = rev(levels(ldaTune$pred$obs)))
plot(ldaROC, legacy.axes = TRUE)
auc(ldaROC)

#### Save results

Test_Error <- append(mean(ldaPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.7882, Test_AUC)
Models <- append("LDA", Models)

```

```

#### PLS DA
set.seed(210)
plsTune <- train(form = Disease~., data = Train2,
                 method = "pls",
                 metric = "ROC",
                 tuneGrid = expand.grid(.ncomp = 1:11),
                 trControl = ctrl)
plsTune

```

```

plsPred <- predict(plsTune, Test2)
table(plsPred, Test2$Disease)

#### Test error rate
mean(plsPred != Test2$Disease)

#### Importance
varImp(plsTune)
plot(varImp(plsTune))

#### Test ROC
library(pROC)

plsROC <- roc(response = plsTune$pred$obs,
              predictor = plsTune$pred$Yes,
              levels = rev(levels(plsTune$pred$obs)))
plot(plsROC, legacy.axes = TRUE)
auc(plsROC)

#### Save results

Test_Error <- append(mean(plsPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.6768, Test_AUC)
Models <- append("PLSDA", Models)

#### Penalized Model
memory.limit(size = 20000)

set.seed(210)
glmnetGrid <- expand.grid(alpha = c(0, .1, .2, .4),
                          lambda = seq(.01, .2, length = 15))

penTune <- train(form = Disease~., data = Train2,
                 method = "glmnet",
                 tuneGrid = glmnetGrid,
                 metric = "ROC",
                 trControl = ctrl)
penTune

penPred <- predict(penTune, Test2)
table(penPred, Test2$Disease)

```

```

#### Test error rate
mean(penPred != Test2$Disease)

#### Importance
varImp(penTune)
plot(varImp(penTune))

#### Test ROC
library(pROC)

penROC <- roc(response = penTune$pred$obs,
               predictor = penTune$pred$Yes,
               levels = rev(levels(penTune$pred$obs)))
plot(penROC, legacy.axes = TRUE)
auc(penROC)

#### Save results

Test_Error <- append(mean(penPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.7833, Test_AUC)
Models <- append("Pen", Models)

#### Nearest Shrunken Centroids
nscGrid <- data.frame(.threshold = 0:25)
nscTune <- train(form = Disease~., data = Train2,
                 method = "pam",
                 tuneGrid = nscGrid,
                 metric = "ROC",
                 trControl = ctrl)
nscTune

nscPred <- predict(nscTune, Test2)
table(nscPred, Test2$Disease)

#### Test error rate
mean(nscPred != Test2$Disease)

#### Importance
varImp(nscTune)
plot(varImp(nscTune))

#### Test ROC
library(pROC)

```

```

nscROC <- roc(response = nscTune$pred$obs,
               predictor = nscTune$pred$Yes,
               levels = rev(levels(nscTune$pred$obs)))
plot(nscROC, legacy.axes = TRUE)
auc(nscROC)

#### Save results

Test_Error <- append(mean(nscPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.6515, Test_AUC)
Models <- append("NSC", Models)

```

```

#### QDA
set.seed(210)
qdaTune <- train(form = Disease~., data = Train2,
                 method = "qda",
                 metric = "ROC",
                 trControl = ctrl)
qdaTune

```

```

qdaPred <- predict(qdaTune, Test2)
table(qdaPred, Test2$Disease)

```

```

#### Test error rate
mean(qdaPred != Test2$Disease)

```

```

#### Importance
varImp(qdaTune)
plot(varImp(qdaTune))

```

```

#### Test ROC
library(pROC)

```

```

qdaROC <- roc(response = qdaTune$pred$obs,
               predictor = qdaTune$pred$Yes,
               levels = rev(levels(qdaTune$pred$obs)))
plot(qdaROC, legacy.axes = TRUE)
auc(qdaROC)

```

```

#### Save results

```

```
Test_Error <- append(mean(qdaPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.7561, Test_AUC)
Models <- append("QDA", Models)
```

```
### MDA
set.seed(210)
mdaTune <- train(form = Disease~., data = Train2,
  method = "mda",
  metric = "ROC",
  tuneGrid = expand.grid(.subclasses = 1:10),
  trControl = ctrl)
mdaTune
```

```
mdaPred <- predict(mdaTune, Test2)
table(mdaPred, Test2$Disease)
```

```
### Test error rate
mean(mdaPred != Test2$Disease)
```

```
### Importance
varImp(mdaTune)
plot(varImp(mdaTune))
```

```
### Test ROC
library(pROC)
```

```
mdaROC <- roc(response = mdaTune$pred$obs,
  predictor = mdaTune$pred$Yes,
  levels = rev(levels(mdaTune$pred$obs)))
plot(mdaROC, legacy.axes = TRUE)
auc(mdaROC)
```

```
### Save results
```

```
Test_Error <- append(mean(mdaPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.7855, Test_AUC)
Models <- append("MDA", Models)
```

```
### FDA
set.seed(210)
```

```

fdaTune <- train(form = Disease~., data = Train2,
                 method = "fda",
                 metric = "ROC",
                 trControl = ctrl)
fdaTune

fdaPred <- predict(fdaTune, Test2)
table(fdaPred, Test2$Disease)

#### Test error rate
mean(fdaPred != Test2$Disease)

#### Importance
varImp(fdaTune)
plot(varImp(fdaTune))

#### Test ROC
library(pROC)

fdaROC <- roc(response = fdaTune$pred$obs,
              predictor = fdaTune$pred$Yes,
              levels = rev(levels(fdaTune$pred$obs)))
plot(fdaROC, legacy.axes = TRUE)
auc(fdaROC)

#### Save results

Test_Error <- append(mean(fdaPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.7785, Test_AUC)
Models <- append("FDA", Models)

#### Naive Bayes
set.seed(210)
nbTune <- train(x = Train2[,c(1:10)], y = Train2$Disease,
               method = "nb",
               metric = "ROC",
               trControl = ctrl)
nbTune

nbPred <- predict(nbTune, Test2)
table(nbPred, Test2$Disease)

#### Test error rate
mean(nbPred != Test2$Disease)

```



```

#### Importance
varImp(nbTune)
plot(varImp(nbTune))

#### Test ROC
library(pROC)

nbROC <- roc(response = nbTune$pred$obs,
             predictor = nbTune$pred$Yes,
             levels = rev(levels(nbTune$pred$obs)))
plot(nbROC, legacy.axes = TRUE)
auc(nbROC)

#### Save results

Test_Error <- append(mean(nbPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.7818, Test_AUC)
Models <- append("NB", Models)

#### KNN
set.seed(210)
knnTune <- train(form = Disease~., data = Train2,
                method = "knn",
                metric = "ROC",
                tuneGrid = data.frame(.k = c(1:100)),
                trControl = ctrl)
knnTune

knnPred <- predict(knnTune, Test2)
table(knnPred, Test2$Disease)

#### Test error rate
mean(knnPred != Test2$Disease)

#### Importance
varImp(knnTune)
plot(varImp(knnTune))

#### Test ROC
library(pROC)

knnROC <- roc(response = knnTune$pred$obs,
             predictor = knnTune$pred$Yes,

```

```

        levels = rev(levels(knnTune$pred$obs)))
plot(knnROC, legacy.axes = TRUE)
auc(knnROC)

#### Save results

Test_Error <- append(mean(knnPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.6518, Test_AUC)
Models <- append("KNN", Models)

```

```

#### Neural Network
set.seed(210)
nnetGrid <- expand.grid(.size = 1:10,
                      .decay = c(0, .01, .1, 0.5))
maxSize <- max(nnetGrid$.size)
numWts <- 200
set.seed(210)
nnetTune <- train(form = Disease~., data = Train2,
                 method = "nnet",
                 metric = "ROC",
                 tuneGrid = nnetGrid,
                 trace = FALSE,
                 maxit = 1000,
                 MaxNWts = numWts,
                 ## ctrl was defined in the previous chapter
                 trControl = ctrl)

```

```

nnetTune

```

```

nnetPred <- predict(nnetTune, Test2)
table(nnetPred, Test2$Disease)

```

```

#### Test error rate
mean(nnetPred != Test2$Disease)

```

```

#### Importance
varImp(nnetTune)
plot(varImp(nnetTune))

```

```

#### Test ROC
library(pROC)

```

```

nnetROC <- roc(response = nnetTune$pred$obs,

```

```

        predictor = nnetTune$pred$Yes,
        levels = rev(levels(nnetTune$pred$obs)))
plot(nnetROC, legacy.axes = TRUE)
auc(nnetROC)

#### Save results

Test_Error <- append(mean(nnetPred != Test2$Disease), Test_Error)
Test_AUC <- append(0.7636, Test_AUC)
Models <- append("NNet", Models)

```

```
##### Plots #####
```

```
#### Make plots comparing fit statistics
```

```
#### Combine into one data frame
```

```
testResults <- data.frame(Test_Error, Test_AUC, Models)
```

```
#### Test error rate, sorted
```

```
OrderErr <- testResults[order(testResults$Test_Error),]
```

```

plot(OrderErr$Test_Err, c(1:11), yaxt = "n",
     main = "Test Error Rate by Model",
     xlab = "Test Error Rate",
     ylab = "Model",
     pch = 16)
axis(2, at = c(1:11), labels = as.character(OrderErr$Models), las = 2,
     cex.axis = 0.8)
text(OrderErr$Test_Err[1:9], c(1:9),
     labels = as.character(round(OrderErr$Test_Err[1:9], 5)),
     adj = -0.2)
text(OrderErr$Test_Err[10:11], c(10:11),
     labels = as.character(round(OrderErr$Test_Err[10:11], 5)),
     adj = 1.2)

```

```

#### AUC, sorted
OrderAUC <- testResults[order(testResults$Test_AUC),]

plot(OrderAUC$Test_AUC, c(1:11), yaxt = "n",
     main = "Test AUC by Model",
     xlab = "Area Under the Curve",
     ylab = "Model",
     pch = 16)
axis(2, at = c(1:11), labels = as.character(OrderAUC$Models), las = 2,
     cex.axis = 0.8)
text(OrderAUC$Test_AUC[1:5], c(1:5),
     labels = as.character(round(OrderAUC$Test_AUC[1:5], 5)),
     adj = -0.2)
text(OrderAUC$Test_AUC[6:11], c(6:11),
     labels = as.character(round(OrderAUC$Test_AUC[6:11], 5)),
     adj = 1.2)

#### Make plot of ROC curves
plot(logicROC, legacy.axis = TRUE, lty = 5, lwd = 3.5, main = "ROC Curves by Model")
lines(ldaROC, col = "red", lty = 3, lwd = 4)
lines(plsROC, col = "blue", lty = 5, lwd = 3.5)
lines(penROC, col = "green", lty = 3, lwd = 4)
lines(nscROC, col = "magenta", lty = 5, lwd = 3.5)
lines(qdaROC, col = "purple", lty = 3, lwd = 4)
lines(mdaROC, col = "orange", lty = 5, lwd = 3.5)
lines(fdaROC, col = "forestgreen", lty = 5, lwd = 3.5)
lines(nbROC, col = "hotpink", lty = 3, lwd = 4)
lines(knnROC, col = "grey50", lty = 3, lwd = 4)
lines(nnetROC, col = "aquamarine", lty = 3, lwd = 4)

legend("topleft",
     legend = as.character(rev(testResults$Models)),
     lty = c(5,3,5,3,5,3,5,3,3,3),
     lwd = c(3.5,4,3.5,4,3.5,4,3.5,3.5,4,4,4),
     col = c("black","red","blue","green","magenta","purple",
             "orange","forestgreen","hotpink","grey50","aquamarine"))

library(lattice)
resamp = resamples(list(Logistic = logicTune, LDA = ldaTune, PLSDA = plsTune,
                       Pen = penTune, NSC = nscTune, QDA = qdaTune,

```

```
MDA = mdaTune, FDA = fdaTune, NB = nbTune,  
KNN = knnTune, NNet = nnetTune))
```

```
dotplot(resamp, metric = "ROC")
```

```
ModelDiff <- diff(resamp)  
ModelDiff$statistics$ROC
```